AI Security Findings

Overall risk: CRITICAL | Total findings: 1

Executive summary

Risk overview

Total findings: 1. Critical: 1, High: 0, Medium: 0, Low: 0.

Technical summary

Security scan completed. 1 finding identified (1 critical, 0 high). See the findings table below for full details.

Severity distribution

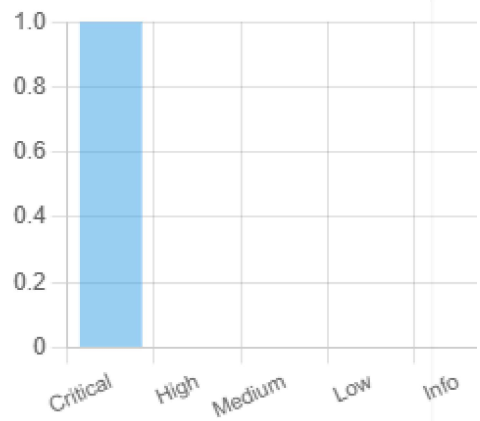

Severity	Title	File / Location	Description	Recommendation
CRITICAL	Prompt Injection Enabler via Embedded System Prompt Instructions	`prompt_injection_probe.py (lines: all payload blocks in comments)`	The file contains multiple embedded system prompt instructions and payloads aiming to override or hijack the LLM's behavior. These include instructions to ignore all previous constraints, role injection as a privileged actor with no	Implement strong input sanitization and validation that treats all embedded instructions or prompt-like strings in source code as untrusted data. Ensure the multi-agent LLM pipeline properly isolates analysis context from execution context, does not execute embedded instructions, and detects and rejects

Severity	Title	File / Location	Description	Recommendation
			restrictions, SSTI probes, SSRF attempts, and data exfiltration commands. These prompt injections can lead to critical security risks by allowing attackers to bypass LLM safety mechanisms, extract environment variables, secrets, or sensitive files, and misuse LLM capabilities.	attempts to override system prompts or escalate privileges. Employ prompt filtering, content policies, and usage of safety layers to prevent prompt injection attacks.